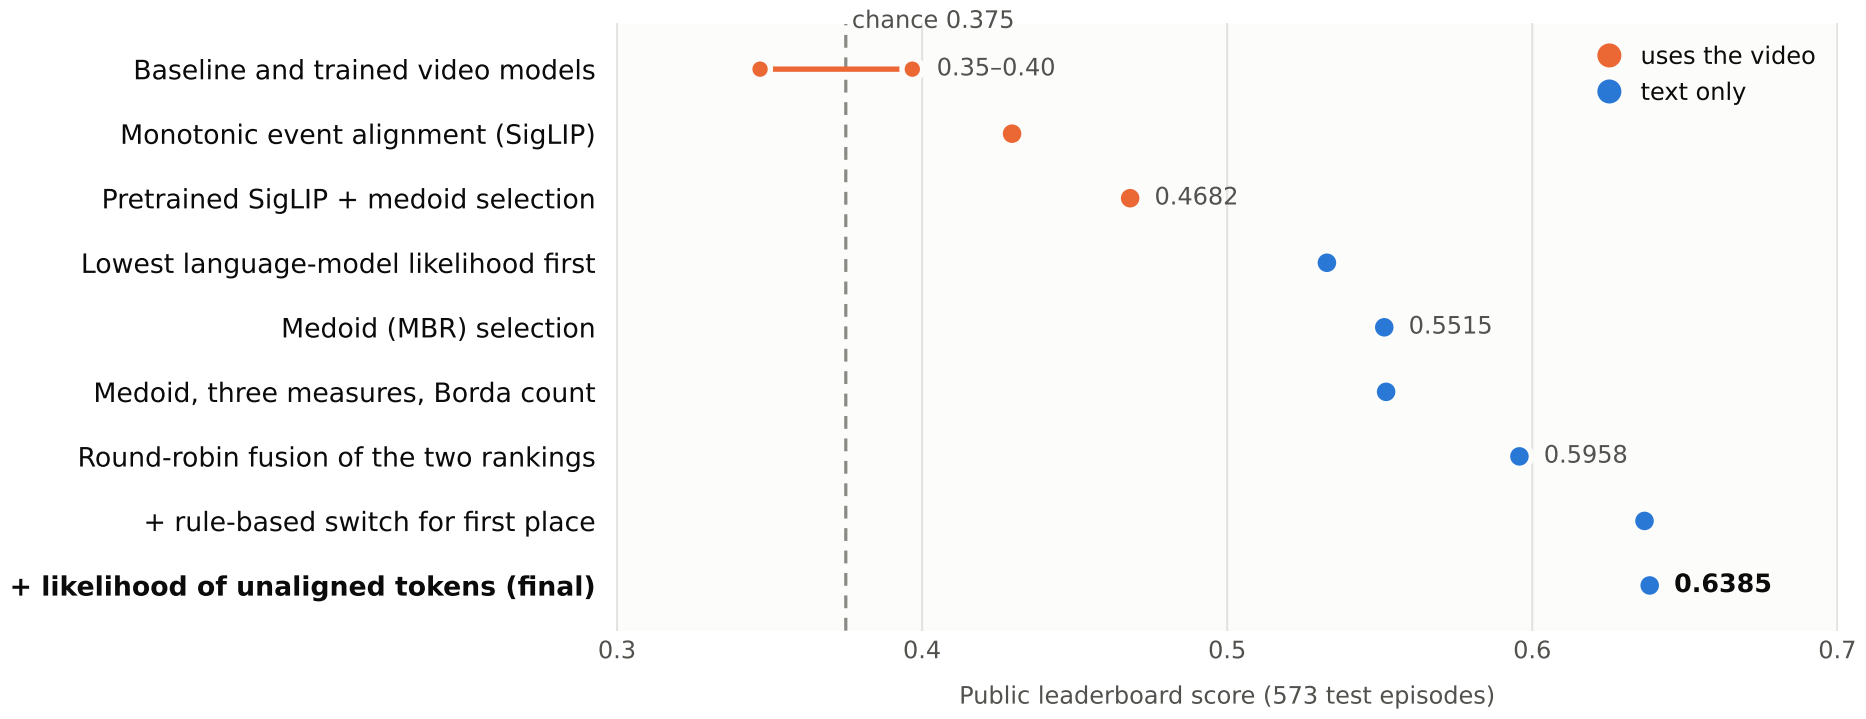
chance 0.375
uses the video
text only
Baseline and trained video models
0.35–0.40
Monotonic event alignment (SigLIP)
Pretrained SigLIP + medoid selection
0.4682
Lowest language-model likelihood first
Medoid (MBR) selection
0.5515
Medoid, three measures, Borda count
Round-robin fusion of the two rankings
0.5958
+ rule-based switch for first place
+ likelihood of unaligned tokens (final)
0.6385
0.3
0.4
0.5
0.6
0.7
Public leaderboard score (573 test episodes)